

CS5710 — Home Assignment 1

Student Name: Manikanth Reddy Devarapalli

Course: Machine Learning, Fall 2025

Department of Computer Science & Cybersecurity, University of Central Missouri

1 — Function Approximation by Hand

Dataset: $(x,y) = \{(1,1), (2,2), (3,2), (4,5)\}$

Model: $\hat{y} = \theta_1 * x + \theta_2$

Mean Squared Error (MSE): $J(\theta) = (1/N) * \sum (y - \hat{y})^2, N = 4$

Model A: $\theta = (1, 0) \rightarrow \hat{y} = x$

Predictions and residuals:

$x=1, y=1 \rightarrow \hat{y}=1$; residual $r = y - \hat{y} = 0$; $r^2 = 0$

$x=2, y=2 \rightarrow \hat{y}=2$; residual $r = y - \hat{y} = 0$; $r^2 = 0$

$x=3, y=2 \rightarrow \hat{y}=3$; residual $r = y - \hat{y} = -1$; $r^2 = 1$

$x=4, y=5 \rightarrow \hat{y}=4$; residual $r = y - \hat{y} = 1$; $r^2 = 1$

Sum of squared residuals = 2; MSE = $0.5 = 2.0/4$

Model B: $\theta = (0.5, 1) \rightarrow \hat{y} = 0.5 * x + 1$

Predictions and residuals:

$x=1, y=1 \rightarrow \hat{y}=1.5$; residual $r = y - \hat{y} = -0.5$; $r^2 = 0.25$

$x=2, y=2 \rightarrow \hat{y}=2$; residual $r = y - \hat{y} = 0$; $r^2 = 0$

$x=3, y=2 \rightarrow \hat{y}=2.5$; residual $r = y - \hat{y} = -0.5$; $r^2 = 0.25$

$x=4, y=5 \rightarrow \hat{y}=3$; residual $r = y - \hat{y} = 2$; $r^2 = 4$

Sum of squared residuals = 4.5; MSE = $1.125 = 4.5/4$

Conclusion: Model A has lower MSE and fits the dataset better.

2 — Random Guessing Practice

Cost function: $J(\theta_1, \theta_2) = 8 * (\theta_1 - 0.3)^2 + 4 * (\theta_2 - 0.7)^2$

Guess $(0.1, 0.2)$:

$(\theta_1 - 0.3) = 0.1 - 0.3 = -0.2 \rightarrow (-0.2)^2 = 0.04 \rightarrow 8 * 0.04 = 0.32$

$(\theta_2 - 0.7) = 0.2 - 0.7 = -0.5 \rightarrow (-0.5)^2 = 0.25 \rightarrow 4 * 0.25 = 1.0$

$J(0.1, 0.2) = 0.32 + 1.0 = 1.32$

Guess (0.5, 0.9):

$$(\theta_1 - 0.3) = 0.5 - 0.3 = 0.2 \rightarrow 0.2^2 = 0.04 \rightarrow 8 * 0.04 = 0.32$$

$$(\theta_2 - 0.7) = 0.9 - 0.7 = 0.2 \rightarrow 0.2^2 = 0.04 \rightarrow 4 * 0.04 = 0.16$$

$$J(0.5, 0.9) = 0.32 + 0.16 = 0.48$$

Which guess is closer to the minimum (0.3, 0.7)?

- $J(0.1, 0.2) = 1.32$; $J(0.5, 0.9) = 0.48$. The smaller value (0.48) corresponds to (0.5, 0.9), so that guess is closer.

Why random guessing is inefficient?

- Random guessing wastes computation and time because it does not follow information about the gradient or the structure of the cost surface.
- Instead, optimization methods (e.g., gradient descent) use local derivative information to move systematically toward minima and reach solutions much faster.

3 — First Gradient Descent Iteration

Dataset: (1,3), (2,4), (3,6), (4,5)

Model: $\hat{y} = \theta_1 * x + \theta_2$; Start: $\theta^*(0) = (0, 0)$; learning rate $\alpha = 0.01$

Step 1: Predictions at $\theta^*(0) = (0, 0)$

For all x , $\hat{y} = 0$. So predictions vector = [0, 0, 0, 0]

Step 2: Residuals $r = y - \hat{y} = y$

$$x=1, y=3, \hat{y}=0 \rightarrow \text{residual } r = 3.0$$

$$x=2, y=4, \hat{y}=0 \rightarrow \text{residual } r = 4.0$$

$$x=3, y=6, \hat{y}=0 \rightarrow \text{residual } r = 6.0$$

$$x=4, y=5, \hat{y}=0 \rightarrow \text{residual } r = 5.0$$

$$\text{Sum } r = 18.0; \text{Sum } x * r = 49.0$$

Step 3: Compute gradient (for MSE $J = (1/N) \sum (y - \hat{y})^2$):

The gradient vector $(\partial J / \partial \theta_1, \partial J / \partial \theta_2) = -(2/N) * [\sum(x * r), \sum(r)]$ where $r = y - \hat{y}$

Therefore gradient at $\theta^*(0) = (-2/4 * 49.0, -2/4 * 18.0) = (-24.5, -9)$

Step 4: Update $\theta^*(1) = \theta^*(0) - \alpha * \text{gradient}$

$$\theta^*(1) = (0, 0) - 0.01 * (-24.5, -9) = (0.245, 0.09)$$

Step 5: Compute $J(\theta^*(0))$ for reference

$$J(\theta^*(0)) = (1/4) * \sum(y^2) = 21.5$$

Now continue from $\theta^*(1)$: compute predictions, residuals, gradient, update to $\theta^*(2)$

Predictions at $\theta^*(1)$:

$$x=1, \hat{y}=0.335, y=3, \text{residual } r = y - \hat{y} = 2.665$$

$$x=2, \hat{y}=0.58, y=4, \text{residual } r = y - \hat{y} = 3.42$$

$$x=3, \hat{y}=0.825, y=6, \text{residual } r = y - \hat{y} = 5.175$$

$x=4, \hat{y}=1.07, y=5$, residual $r = y - \hat{y} = 3.93$
 Sum r (at $\theta^{(1)}$) = 15.19; Sum $x*r$ (at $\theta^{(1)}$) = 40.75
 Gradient at $\theta^{(1)} = (-20.375, -7.595)$
 $\theta^{(2)} = \theta^{(1)} - \alpha * \text{gradient} = (0.44875, 0.16595)$

Compare training loss:

$J(\theta^{(0)}) = 21.5$
 $J(\theta^{(1)}) = 15.256$
 $J(\theta^{(2)}) = 10.9223$

4 — Compare Random Guessing vs Gradient Descent

Dataset: (1,2), (2,2), (3,4), (4,6)

Two random guesses and one-step gradient descent from $\theta=(0,0)$ with $\alpha=0.01$

Random guess A: $\theta = (0.2, 0.5)$

Compute predictions $\hat{y}_i = 0.2 * x_i + 0.5$, residuals and squared residuals; MSE = 8.35

Random guess B: $\theta = (0.9, 0.1)$

MSE = 1.935

Gradient Descent one-step from $\theta=(0,0)$:

radient computed from data: -21, -7

θ after one update: (0.21, 0.07)

MSE at this $\theta = 10.5091$

Conclusion: Compare the numeric MSEs above. Gradient descent should produce a θ with lower MSE than most random guesses because it moves in the direction of steepest descent reducing the cost.

5 — Recognizing Underfitting and Overfitting

If training error is very high and test error is also very high → Underfitting.

Why: The model lacks capacity (too simple) to capture patterns in the training data. High training error shows the model can't fit the training set.

Fixes (pick two):

1) Increase model complexity (use more features, polynomial terms, or a more flexible model).

2) Reduce regularization (if using strong regularization like large λ in ridge), or train longer with better hyperparameters.

Other options: perform feature engineering, add more relevant features, or use a model with lower bias.

6 — Comparing Models

Given: Model A fits training data almost perfectly but performs poorly on unseen data.

Model B does not fit training data well and performs poorly on unseen data.

Model A: Overfitting — low bias, high variance. The model memorizes noise/peculiarities in the training set and fails to generalize.

Remedies: add regularization (L2/L1), reduce model complexity, get more training data, or use cross-validation to tune hyperparameters.

Model B: Underfitting — high bias, low variance. The model is too simple to capture true relationships.

Remedies: increase model complexity (e.g., more features, polynomial features), reduce regularization, or try a different model family.

Bias-Variance tradeoff summary: Increasing complexity reduces bias but raises variance; decreasing complexity increases bias but lowers variance. Choose model complexity to balance both.